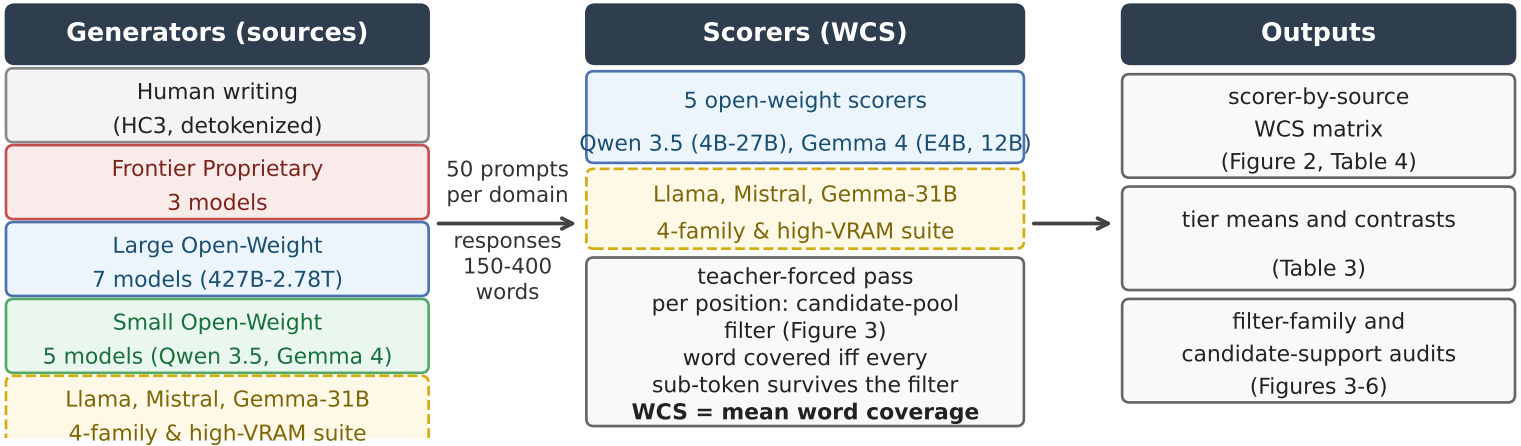
Evaluation pipeline: every source response is scored by every scorer with the Word Coverage Score (WCS)
Generators (sources)
Human writing
(HC3, detokenized)
Frontier Proprietary
3 models
Large Open-Weight
7 models (427B-2.78T)
Small Open-Weight
5 models (Qwen 3.5, Gemma 4)
Llama, Mistral, Gemma-31B
4-family & high-VRAM suite
50 prompts
per domain
responses
150-400
words
Scorers (WCS)
5 open-weight scorers
Qwen 3.5 (4B-27B), Gemma 4 (E4B, 12B)
Llama, Mistral, Gemma-31B
4-family & high-VRAM suite
teacher-forced pass
per position: candidate-pool
filter (Figure 3)
word covered iff every
sub-token survives the filter
WCS = mean word coverage
Outputs
scorer-by-source
WCS matrix
(Figure 2, Table 4)
tier means and contrasts
(Table 3)
filter-family and
candidate-support audits
(Figures 3-6)